

## Computer Science & Information Systems

# Big Data Systems – Lab Sheet

## HBASE

### Objectives

Students should be able to

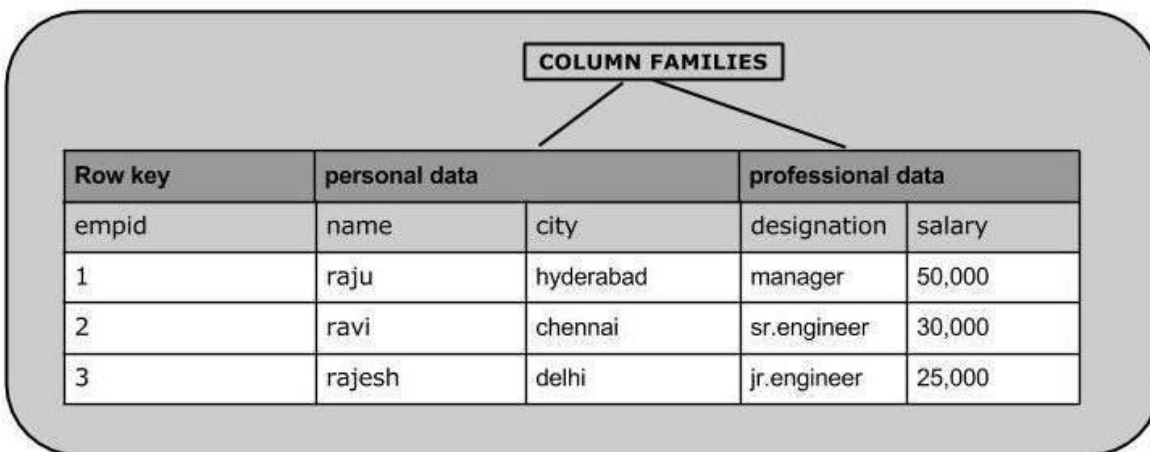
- A. Gain understanding about HBASE
- B. Store, manipulate and retrieve data using HBASE queries

### Introduction to HBASE

HBASE is column oriented non-relational database management system that provides capability to access large files in HDFS. It is a key-value store with no fixed schema unlike RDBMS. It is typically used for analytical queries that access specific columns. It is strongly consistent data store.

### Columnar Storage

Data in a row is collection of column families with each column being key-value pair.



The advantage of having creating column families is that queries that access information in a particular column family will run faster. For example, queries that access only personal data or only professional data in the above snapshot will run faster.

## HBASE QUERIES

In this section we will discuss various clauses/operators used for HBASE queries.

In order to run HBASE queries, we need to start HMaster as follows.

```
[centos@master~]cd /opt/hbase-2.4.15/bin  
[centos@master bin]$ ./start-hbase.sh
```

You can check if HMaster is running by executing jps as follows

```
[centos@master bin]$ jps
```

```
[centos@master bin]$ jps  
802 QuorumPeerMain  
9410 HMaster  
794 Master  
2170 NameNode  
2298 DataNode  
4058 NodeManager  
3179 SecondaryNameNode  
3931 ResourceManager  
9883 Jps
```

In order to access HBASE shell use the following command

```
[centos@master bin]$ hbase shell
```

## CREATE

The create command is used to create a table. While creating a table you must specify the table name and the Column Family name. The syntax to create a table in HBase shell is shown below.

### Syntax

```
hbase> create '<table-name>', '<column-family-name>'
```

### Example

```
hbase> create 'emp' , 'personal information' , 'professional  
information'
```

The above command will create a table with emp having 2 column families 'personal information' and 'professional information'. You need to specify at least one column family in order to create a table.

### Output

```
File Edit View Search Terminal Help  
hbase:001:0> create 'emp', 'personal information', 'professional information'  
Created table emp  
Took 1.1337 seconds  
=> Hbase::Table - emp
```

## LIST TABLES

The list command is used to list all the tables;

```
hbase> list;
```

The above command will list all the tables as follows.

```
File Edit View Search Terminal Help  
hbase:004:0> list  
TABLE  
emp  
1 row(s)  
Took 0.0074 seconds  
=> ["emp"]
```

## DISABLING A TABLE

To delete a table or change its settings, you need to first disable the table using the disable command. The syntax is as follows.

### Syntax

```
hbase> disable '<table-name>'
```

#### Example

```
hbase> disable 'emp'
```

The above command will disable the table 'emp'.

#### Output

```
File Edit View Search Terminal Help
hbase:005:0> disable 'emp'
Took 0.4365 seconds
hbase:006:0>
```

#### DISABLE\_ALL

This command is used to disable all the tables matching the given regex. The syntax for disable\_all command is given below.

##### Syntax

```
hbase>disable_all '<regex>'
```

#### Example

```
hbase>disable_all 'e.*'
```

The above command will disable all the tables starting with e.

#### Output

```
File Edit View Search Terminal Help
hbase:012:0> disable_all 'e.*'
emp

Disable the above 1 tables (y/n)?
y
1 tables successfully disabled
Took 3.1088 seconds
```

## ENABLING A TABLE

You need to enable a disabled table before you can use it. The syntax is as follows.

### Syntax

```
hbase> enable '<table-name>'
```

### Example

```
hbase> enable 'emp'
```

The above command will enable the table 'emp'.

### Output

```
File Edit View Search Terminal Help
hbase:013:0> enable 'emp'
Took 0.6365 seconds
```

## DESCRIBE

The DESCRIBE command returns the description of the table. The syntax is as follows

### Syntax

```
hbase> describe '<table-name>'
```

### Example

```
hbase> describe 'emp'
```

### Output

```
File Edit View Search Terminal Help
hbase:014:0> describe 'emp'
Table emp is ENABLED
emp
COLUMN FAMILIES DESCRIPTION
{NAME => 'personal information', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', COMPRESSION => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
{NAME => 'professional information', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', COMPRESSION => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
2 row(s)
Quota is disabled
Took 0.0330 seconds
```

## ALTER

The ALTER command can be used to make changes to existing table. Using this command, you can change the maximum number of cells of a column family, set and delete table scope operators, and delete a column family from a table.

### Use of alter to change maximum number of cells of a column family

#### Syntax

```
hbase> alter '<table-name>', NAME=> '<cf-name>', VERSIONS => 5
```

#### Example

```
hbase> alter 'emp', NAME=> 'personal information', VERSIONS => 5
```

#### Output

```
File Edit View Search Terminal Help
hbase:015:0> alter 'emp', NAME=> 'personal information', VERSIONS=>5
Updating all regions with the new schema...
1/1 regions updated.
Done.
Took 1.8403 seconds
```

### Use of alter to set table scope operators

Using alter, you can set table scope operators such as MAX\_FILESIZE, READONLY, MEMSTORE\_FLUSH\_SIZE, DEFERRED\_LOG\_FLUSH, etc.

#### Syntax

```
hbase> alter '<table-name>', <table-scope-operator>
```

#### Example

```
hbase> alter 'emp', READONLY
```

#### Output

```
File Edit View Search Terminal Help
hbase:017:0> alter 'emp', READONLY
Updating all regions with the new schema...
1/1 regions updated.
Done.
Took 1.6346 seconds
```

### Use of alter to remove table scope operators

Using alter, you can also remove table scope operators such as MAX\_FILESIZE, READONLY, MEMSTORE\_FLUSH\_SIZE, DEFERRED\_LOG\_FLUSH, etc.

#### Syntax

```
hbase> alter '<table-name>', <table-scope-operator>
```

#### Example

```
hbase> alter 'emp', METHOD => 'table_att_unset',  
NAME=>READONLY
```

### Use of alter to remove a column family

Using alter, you can also delete a column family.

#### Syntax

```
hbase> alter '<table-name>', delete=> 'cf-name'
```

#### Example

```
hbase> alter 'emp', delete=> 'professional information'
```

### EXISTS

You can verify the existence of a table using the exists command. The following example shows how to use this command.

#### Syntax

```
hbase> exists '<table-name>'
```

#### Example

```
hbase> exists 'emp'
```

File	Edit	View	Search	Terminal	Help
hbase:005:0> exists 'emp'					
Table emp does exist					
Took 0.0074 seconds					
=> true					

## DROP A TABLE

Using the drop command, you can delete a table. Before dropping a table, you have to disable it.

### Syntax

```
hbase> disable '<table-name>'  
hbase> drop '<table-name>'
```

### Example

```
hbase> disable 'emp'  
hbase> drop 'emp'
```

## Output

```
File Edit View Search Terminal Help  
hbase:006:0> disable 'emp'  
Took 0.3541 seconds  
hbase:007:0> drop 'emp'  
Took 0.1330 seconds  
hbase:008:0> list  
TABLE  
0 row(s)  
Took 0.0162 seconds  
=> []
```

## DROP ALL

The drop\_all command is used to drop the tables matching the “regex” given in the command. Before dropping a table, you must disable it. Its syntax is as follows:

### Syntax

```
hbase>disable_all '<regex>'  
hbase>drop_all '<regex>'
```

### Example

```
hbase>disable_all 'e.*'  
hbase>drop_all 'e.*'
```



The above command will first disable all the tables starting with e and then drop all those tables.

```
File Edit View Search Terminal Help
hbase:003:0> disable_all 'e.*'
emp

Disable the above 1 tables (y/n)?
y
1 tables successfully disabled
Took 3.5940 seconds
hbase:004:0> drop_all 'e.*'
emp

Drop the above 1 tables (y/n)?
y
1 tables successfully dropped
Took 4.1493 seconds
hbase:005:0> list
TABLE
0 row(s)
Took 0.0085 seconds
=> []
```

## INSERTING DATA INTO TABLES

Using **put** command, you can insert rows into a table. Its syntax is as follows:

### Syntax

```
hbase> put '<tab-name>', 'row-num', '<colfamily:colname>',
'<value>'
```

### Example

```
hbase> put 'emp', '1', 'personal information:name', 'raju'
```

## SCAN

The scan command is used to view the data in HBASE table.

### Syntax

```
hbase> scan '<table-name>'
```

### Example

```
hbase> scan 'emp'
```

## Output

```
File Edit View Search Terminal Help
hbase:004:0> put 'emp', '1', 'personal information:name', 'raju'
Took 0.2997 seconds
hbase:005:0> scan 'emp'
ROW                                COLUMN+CELL
1                                  column=personal information:name, timestamp=2023-01-05T19:13:25.560, value=raju
1 row(s)
Took 0.0486 seconds
```

## UPDATING DATA IN TABLES

Using **put** command, you can also update data stored in a table. Its syntax is as follows:

### Syntax

```
hbase> put '<tab-name>', 'row-num', '<colfamily:colname>',  
'<new-value>'
```

### Example

```
hbase> put 'emp' '1' 'personal information:name', 'raj'
```

### Output

```
File Edit View Search Terminal Help
hbase:030:0> put 'emp', '1', 'personal information:name','raj'
Took 0.0072 seconds
hbase:031:0> scan 'emp'
ROW                                COLUMN+CELL
1                                  column=personal information:name, timestamp=2023-01-05T19:18:54.838, value=raj
1 row(s)
Took 0.0077 seconds
```

## READING DATA FROM TABLES

Using **get** command, you can read the data from a table. Its syntax is as follows:

### Syntax: Reading a specific row

```
hbase> get '<tab-name>', 'row-num'
```

### Example

```
hbase> get 'emp', '1'
```

### Output

```
File Edit View Search Terminal Help
hbase:032:0> get 'emp', 1
COLUMN                                CELL
personal information:name             timestamp=2023-01-05T19:18:54.838, value=raj
1 row(s)
Took 0.0374 seconds
```

Syntax: Reading a specific column

```
hbase> get '<tab-name>', 'row-num' {COLUMN=> 'cf-name:col-  
name' }
```

Example

```
hbase> get 'emp' '1' {COLUMN => 'personalinformation:name' }
```

**Output**

```
File Edit View Search Terminal Help  
hbase:002:0> get 'emp',1,{COLUMN => 'personal information:'}  
COLUMN                                CELL  
personal information:name              timestamp=2023-01-05T19:18:54.838, value=raj  
1 row(s)  
Took 0.0251 seconds
```

## DELETE DATA FROM TABLES

Using **delete** command, you can delete specific cells in a table. Its syntax is as follows:

Syntax:

```
hbase> delete '<tab-name>', 'row-num', '<col-name>', '<time-  
stamp>'
```

Example

```
hbase>delete,'emp' 1, 'personal information:name'
```

**Output**

```
File Edit View Search Terminal Help  
hbase:008:0> delete 'emp',1,'personal information:name'  
Took 0.0068 seconds  
hbase:009:0> scan 'emp'  
ROW                                COLUMN+CELL  
0 row(s)  
Took 0.0050 seconds
```

Using the “deleteall” command, you can delete all the cells in a row.

### Syntax:

```
hbase>deleteall '<tab-name>','row-num'
```

### Example

```
hbase>deleteall 'emp' '1'
```

### Output

```
File Edit View Search Terminal Help
hbase:012:0> scan 'emp'
ROW                                COLUMN+CELL
1                                  column=personal information:address, timestamp=2023-01-05T19:45:03.237, value=delhi
1                                  column=personal information:name, timestamp=2023-01-05T19:43:34.635, value=raju
1                                  column=professional information:designation, timestamp=2023-01-05T19:46:34.111, value=SDE
1                                  column=professional information:salary, timestamp=2023-01-05T19:47:27.642, value=100000
2                                  column=personal information:address, timestamp=2023-01-05T19:48:57.076, value=pune
2                                  column=personal information:name, timestamp=2023-01-05T19:48:25.322, value=seema
2                                  column=professional information:designation, timestamp=2023-01-05T19:49:57.521, value=BA
2                                  column=professional information:salary, timestamp=2023-01-05T19:50:41.805, value=150000
2 row(s)
Took 0.0183 seconds
hbase:013:0> deleteall 'emp',1
Took 0.0093 seconds
hbase:014:0> scan 'emp'
ROW                                COLUMN+CELL
2                                  column=personal information:address, timestamp=2023-01-05T19:48:57.076, value=pune
2                                  column=personal information:name, timestamp=2023-01-05T19:48:25.322, value=seema
2                                  column=professional information:designation, timestamp=2023-01-05T19:49:57.521, value=BA
2                                  column=professional information:salary, timestamp=2023-01-05T19:50:41.805, value=150000
1 row(s)
Took 0.0107 seconds
```

## COUNT

The count command is used to count number of rows in HBASE table.

### Syntax

```
hbase> count '<table-name>'
```

### Example

```
hbase> count 'emp'
```

### Output

```
File Edit View Search Terminal Help
hbase:020:0> scan 'emp'
ROW                                COLUMN+CELL
1                                  column=personal information:address, timestamp=2023-01-05T19:55:27.330, value=delhi
1                                  column=personal information:name, timestamp=2023-01-05T19:55:14.601, value=raju
1                                  column=professional information:designation, timestamp=2023-01-05T19:54:53.228, value=SDE
1                                  column=professional information:salary, timestamp=2023-01-05T19:54:44.597, value=100000
2                                  column=personal information:address, timestamp=2023-01-05T19:48:57.076, value=pune
2                                  column=personal information:name, timestamp=2023-01-05T19:48:25.322, value=seema
2                                  column=professional information:designation, timestamp=2023-01-05T19:49:57.521, value=BA
2                                  column=professional information:salary, timestamp=2023-01-05T19:50:41.805, value=150000
2 row(s)
Took 0.0111 seconds
hbase:021:0> count 'emp'
2 row(s)
Took 0.0425 seconds
=> 2
```

## TRUNCATE

The truncate command disables drop and recreates a table

### Syntax

```
hbase> truncate '<table-name>'
```

### Example

```
hbase> truncate 'emp'
```

### Output

```
File Edit View Search Terminal Help
hbase:022:0> scan 'emp'
ROW                                COLUMN+CELL
1                                  column=personal information:address, timestamp=2023-01-05T19:55:27.330, value=delhi
1                                  column=personal information:name, timestamp=2023-01-05T19:55:14.601, value=raju
1                                  column=professional information:designation, timestamp=2023-01-05T19:54:53.228, value=SDE
1                                  column=professional information:salary, timestamp=2023-01-05T19:54:44.597, value=100000
2                                  column=personal information:address, timestamp=2023-01-05T19:48:57.076, value=pune
2                                  column=personal information:name, timestamp=2023-01-05T19:48:25.322, value=seema
2                                  column=professional information:designation, timestamp=2023-01-05T19:49:57.521, value=BA
2                                  column=professional information:salary, timestamp=2023-01-05T19:50:41.805, value=150000
2 row(s)
Took 0.0114 seconds
hbase:023:0> truncate 'emp'
Truncating 'emp' table (it may take a while):
Disabling table...
Truncating table...
Took 1.9005 seconds
hbase:024:0> scan 'emp'
ROW                                COLUMN+CELL
0 row(s)
Took 0.8300 seconds
```

## SingleColumnValueFilter

In order to filter the rows on the HBase shell using Scan, you need to import the org.apache.hadoop.hbase.filter.SingleColumnValueFilter class along with some other class explained below

### Syntax

```
hbase>import org.apache.hadoop.hbase.filter.SingleColumnValueFilter
hbase> import org.apache.hadoop.hbase.filter.CompareFilter
hbase> import org.apache.hadoop.hbase.filter.BinaryComparator
```

```
File Edit View Search Terminal Help
hbase:023:0> import org.apache.hadoop.hbase.filter.SingleColumnValueFilter
=> [Java::OrgApacheHadoopHbaseFilter::SingleColumnValueFilter]
hbase:024:0> import org.apache.hadoop.hbase.filter.CompareFilter
=> [Java::OrgApacheHadoopHbaseFilter::CompareFilter]
hbase:025:0> import org.apache.hadoop.hbase.filter.BinaryComparator
=> [Java::OrgApacheHadoopHbaseFilter::BinaryComparator]
```

### Example-1: Comparison with name

```
hbase>scan 'emp', { FILTER
=>SingleColumnValueFilter.new(Bytes.toBytes('personal information'),
Bytes.toBytes('name'),
CompareFilter::CompareOp.valueOf('EQUAL'),BinaryComparator.new(Bytes
.toBytes('seema')) ) }
```

The above query will return the details corresponding to name 'seema' as follows

```
File Edit View Search Terminal Help
hbase:026:0> scan 'emp'
COLUMN+CELL
1 column=personal information:address, timestamp=2023-01-06T20:30:35.505, value=delhi
1 column=personal information:name, timestamp=2023-01-06T20:30:19.535, value=raju
1 column=professional information:designation, timestamp=2023-01-06T20:34:30.741, value=SDE
1 column=professional information:salary, timestamp=2023-01-06T20:34:50.828, value=100000
2 column=personal information:address, timestamp=2023-01-06T20:31:25.658, value=pune
2 column=personal information:name, timestamp=2023-01-06T20:31:04.204, value=seema
2 column=professional information:designation, timestamp=2023-01-06T20:35:49.402, value=BA
2 column=professional information:salary, timestamp=2023-01-06T20:36:32.448, value=150000
2 row(s)
Took 0.0271 seconds
hbase:027:0> scan 'emp', { FILTER => SingleColumnValueFilter.new(Bytes.toBytes('personal information'), Bytes.toBytes('name'), CompareFilter::CompareOp.valueOf('EQUAL'), BinaryComparator.new(Bytes.toBytes('seema')) ) }
COLUMN+CELL
2 column=personal information:address, timestamp=2023-01-06T20:31:25.658, value=pune
2 column=personal information:name, timestamp=2023-01-06T20:31:04.204, value=seema
2 column=professional information:designation, timestamp=2023-01-06T20:35:49.402, value=BA
2 column=professional information:salary, timestamp=2023-01-06T20:36:32.448, value=150000
1 row(s)
Took 0.0739 seconds
```

### Example-2: Comparison with salary

```
hbase>scan 'emp', { FILTER
=>SingleColumnValueFilter.new(Bytes.toBytes('professional
information'), Bytes.toBytes('salary'),
CompareFilter::CompareOp.valueOf('GREATER'),BinaryComparator.new(Bytes
.toBytes('100000')) ) }
```

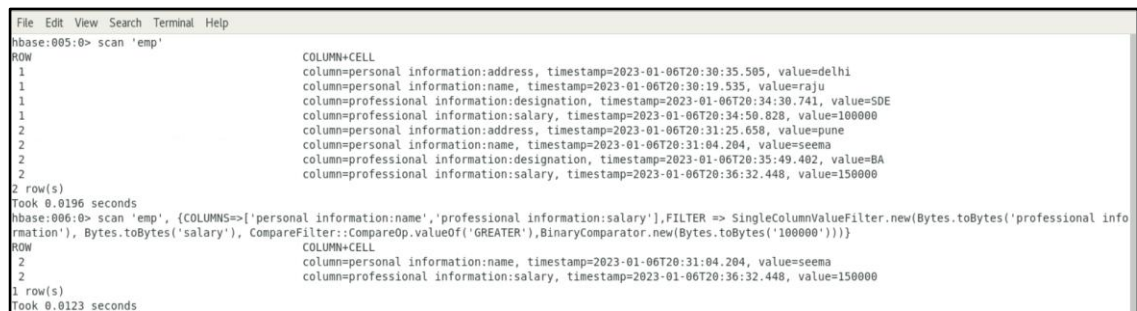
The above query will return the rows where salary>100000 as follows

```
File Edit View Search Terminal Help
hbase:036:0> scan 'emp'
COLUMN+CELL
1 column=personal information:address, timestamp=2023-01-06T20:30:35.505, value=delhi
1 column=personal information:name, timestamp=2023-01-06T20:30:19.535, value=raju
1 column=professional information:designation, timestamp=2023-01-06T20:34:30.741, value=SDE
1 column=professional information:salary, timestamp=2023-01-06T20:34:50.828, value=100000
2 column=personal information:address, timestamp=2023-01-06T20:31:25.658, value=pune
2 column=personal information:name, timestamp=2023-01-06T20:31:04.204, value=seema
2 column=professional information:designation, timestamp=2023-01-06T20:35:49.402, value=BA
2 column=professional information:salary, timestamp=2023-01-06T20:36:32.448, value=150000
2 row(s)
Took 0.0137 seconds
hbase:037:0> scan 'emp', { FILTER => SingleColumnValueFilter.new(Bytes.toBytes('professional information'), Bytes.toBytes('salary'), CompareFilter::CompareOp.valueOf('GREATER'), BinaryComparator.new(Bytes.toBytes('100000')) ) }
COLUMN+CELL
2 column=personal information:address, timestamp=2023-01-06T20:31:25.658, value=pune
2 column=personal information:name, timestamp=2023-01-06T20:31:04.204, value=seema
2 column=professional information:designation, timestamp=2023-01-06T20:35:49.402, value=BA
2 column=professional information:salary, timestamp=2023-01-06T20:36:32.448, value=150000
1 row(s)
Took 0.0099 seconds
```

### Example-3: Displaying selected columns based on filtering condition

```
hbase>scan 'emp', {COLUMNS=>['personal
information:name','professional information:salary'],FILTER
=>SingleColumnValueFilter.new(Bytes.toBytes('professional
information'), Bytes.toBytes('salary'),
CompareFilter::CompareOp.valueOf('GREATER'),BinaryComparator.new(Bytes
.toBytes('100000')))}
```

The above query will return the name and salary where salary>100000



```
File Edit View Search Terminal Help
hbase:005:0> scan 'emp'
ROW COLUMN+CELL
1 column=personal information:address, timestamp=2023-01-06T20:30:35.505, value=delhi
1 column=personal information:name, timestamp=2023-01-06T20:30:19.535, value=raju
1 column=professional information:designation, timestamp=2023-01-06T20:34:30.741, value=SDE
1 column=professional information:salary, timestamp=2023-01-06T20:34:50.828, value=100000
2 column=personal information:address, timestamp=2023-01-06T20:31:25.658, value=pune
2 column=personal information:name, timestamp=2023-01-06T20:31:04.204, value=seema
2 column=professional information:designation, timestamp=2023-01-06T20:35:49.402, value=BA
2 column=professional information:salary, timestamp=2023-01-06T20:36:32.448, value=150000
2 row(s)
Took 0.0196 seconds
hbase:006:0> scan 'emp', {COLUMNS=>['personal information:name','professional information:salary'],FILTER => SingleColumnValueFilter.new(Bytes.toBytes('professional info
rmation'), Bytes.toBytes('salary'), CompareFilter::CompareOp.valueOf('GREATER'),BinaryComparator.new(Bytes.toBytes('100000')))}
ROW COLUMN+CELL
2 column=personal information:name, timestamp=2023-01-06T20:31:04.204, value=seema
2 column=professional information:salary, timestamp=2023-01-06T20:36:32.448, value=150000
1 row(s)
Took 0.0123 seconds
```

## Outputs/Results

- Students should be able to appreciate the usage of HBASE.
- Students should be able to appreciate column families of HBASE
- Students should be able to execute queries using various options available

## Observations

Students should carefully observe the syntax of HBASE queries and the output



## References

- [Tutorial point](#)
- [Spark By Examples](#)