SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

SCSI – The Protocol for all Storage Architectures

David Deming, Solution Technology

April 12, 2005

Abstract

SCSI: The Protocol for All Storage Architectures

This session will appeal to System Administrators, Storage Administrators, Storage Architects, and those that are seeking a fundamental understanding of SCSI Protocol and how it benefits your IT storage applications. The session will delve into the SCSI model, it's protocol, and how storage applications benefit from having a single high level protocol. The audience will receive the fundamental understanding of why SCSI is used as the storage industry's main language.

SNIA Legal Notice

- The material contained in this tutorial is copyrighted by the SNIA.
- Member companies and individuals may use this material in presentations and literature under the following conditions:
  - Any slide or slides used must be reproduced without modification
  - The SNIA must be acknowledged as source of any material used in the body of any document containing material from these presentations.
- This presentation is a project of the SNIA Education Committee.

Objectives

- Brief understanding of SCSI and its history
- Understanding of industry standards and the role of the SCSI Architecture
- Understanding the SCSI Architecture model
- Understanding the I/O Operation and the importance behind SCSI Protocol
- Learn how parallel SCSI, Fibre Channel, and iSCSI rely on the SCSI Architecture to support storage applications

Brief History

- SCSI-1: In the beginning – 1987
  - **SCSI was adapted from the SASI Interface (1979)**
  - **Specified a physical interface, transport protocol, and standard command for disk, tape, and other storage devices**
- SCSI-2: Second coming of SCSI
  - **Speed and bus width enhancements**
  - **Multi-threading commands and transport protocol cleanup**
  - **Added more storage device type command sets**
- SCSI-3: Today and beyond
  - **More speed: Ultra-2/Ultra-160/Ultra-320**
  - **Separation of Physical Interface, Transport Protocols, and the SCSI Command Sets**

SCSI Today and the Future

- SCSI Language
  - **Every architecture uses the SCSI command sets**
  - **Every architecture uses the SCSI upper layer protocol model**
  - **SCSI is the Language of Love**
- SCSI Transport
  - **Defines protocol mapping and transfer of SCSI language**
  - **Fibre Channel Protocol (FCP), Serial Storage Architecture (SSA), Serial Bus Protocol (SBP), iSCSI, Serial ATA (SATA), Serial Attached SCSI (SAS)**
- SCSI Physical Architectures
  - **Serial Attached SCSI – next generation SCSI physical interface**
  - **Fibre Channel, SSA, 1394, ATAPI, TCP/IP**

Introduction to Standards

SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

SCSI Standards and their impact on all storage interfaces

Standards Organizations

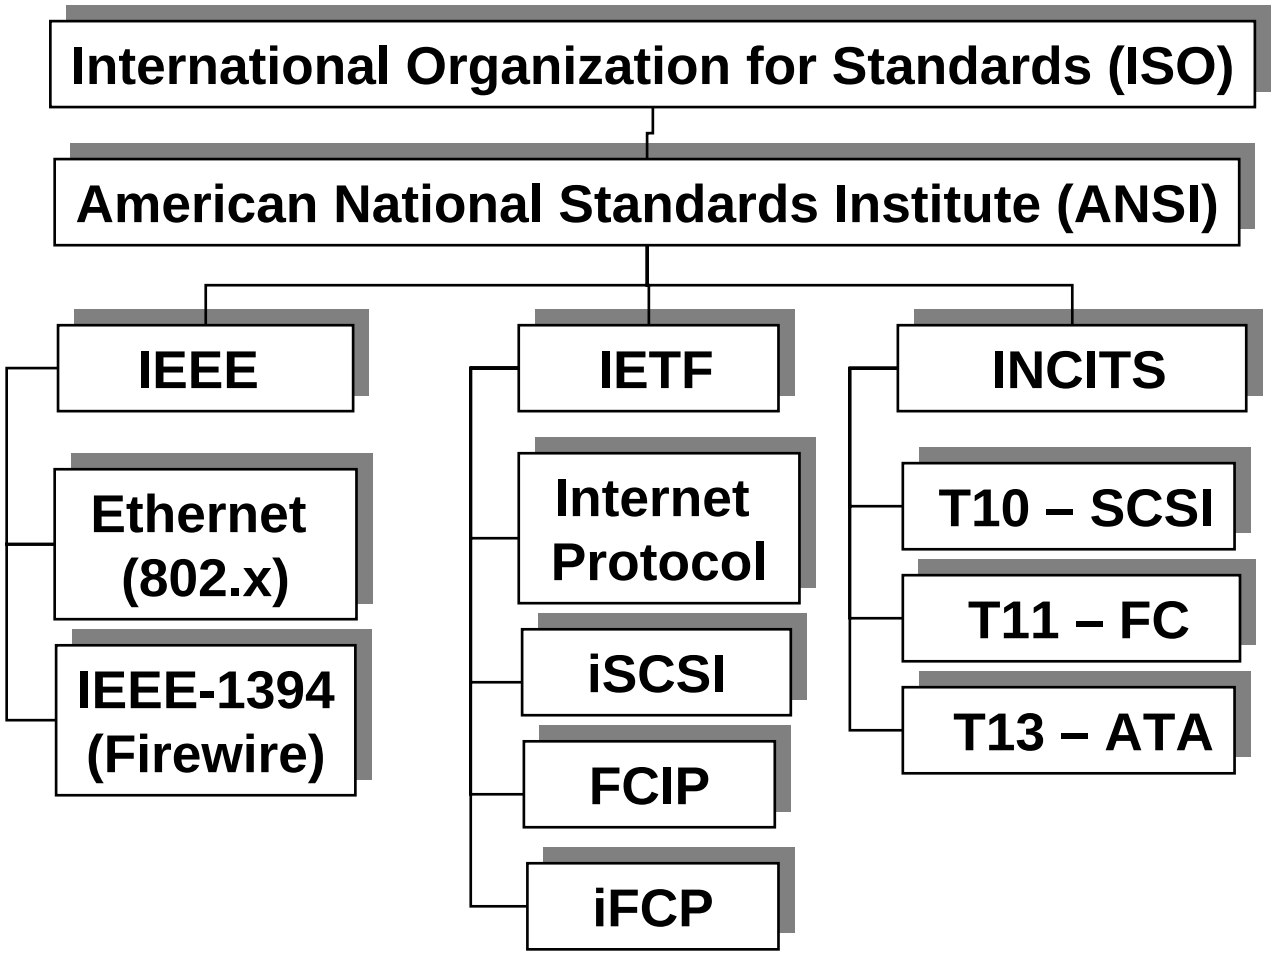

SCSI Standards

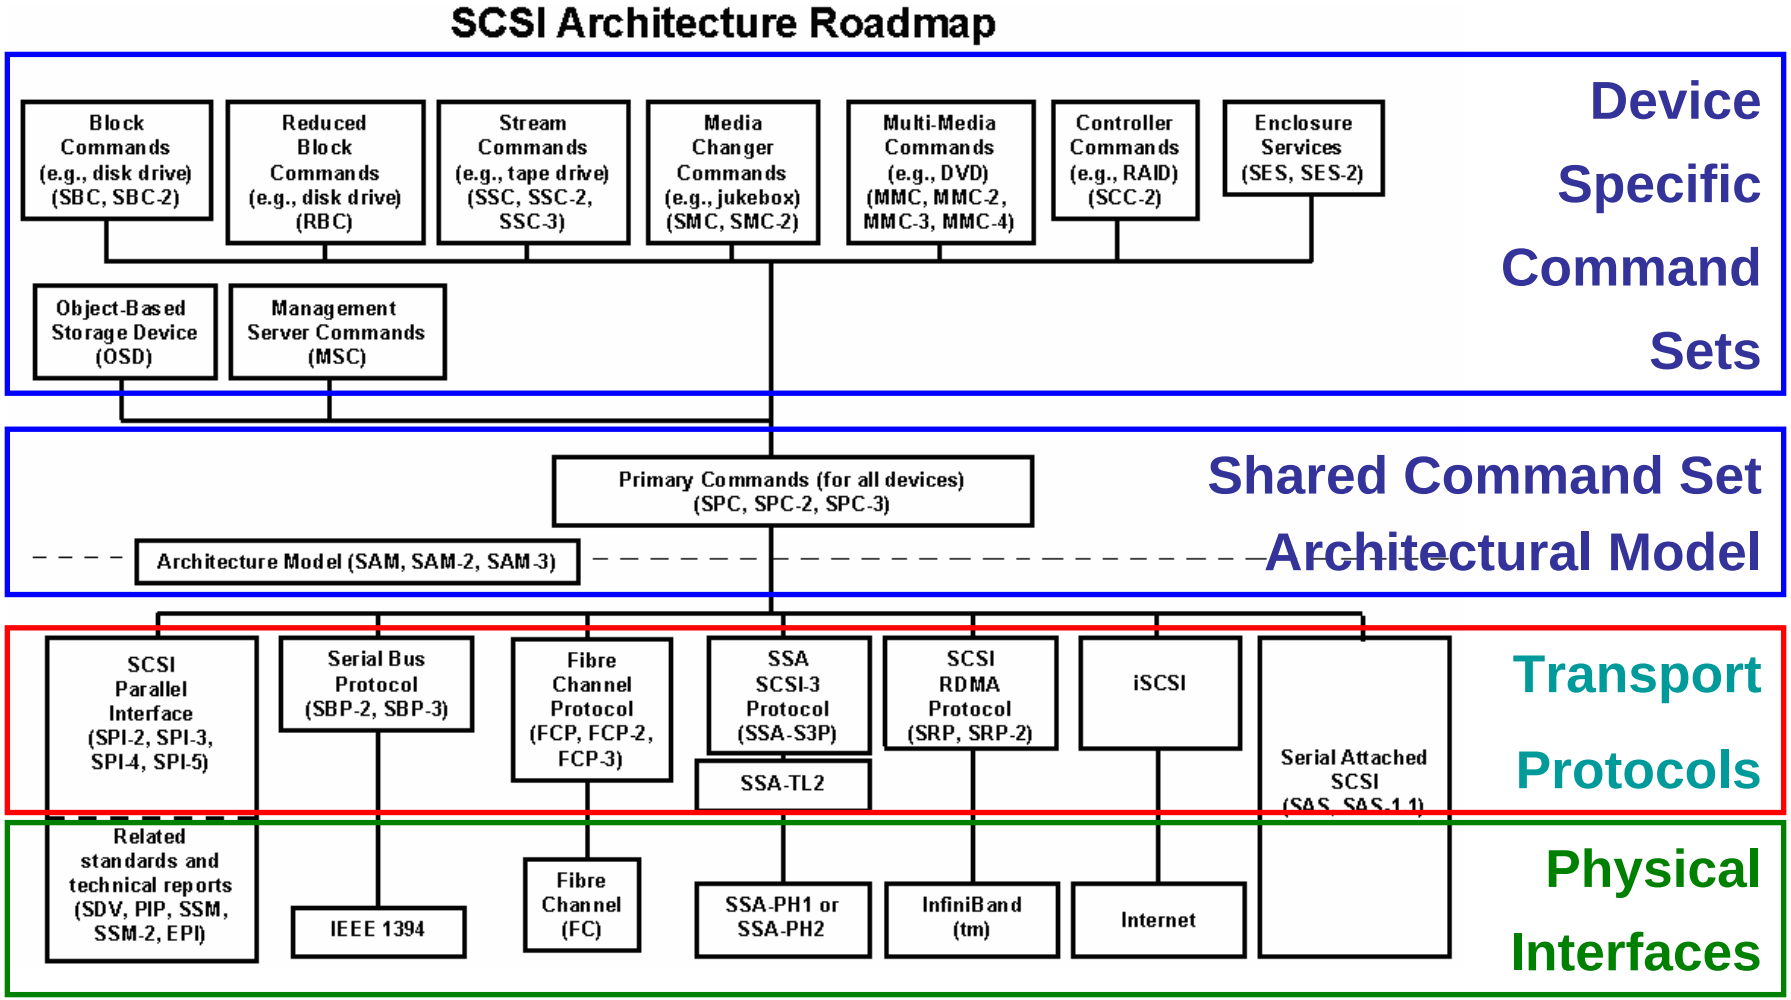

SCSI-3 Physical Interfaces

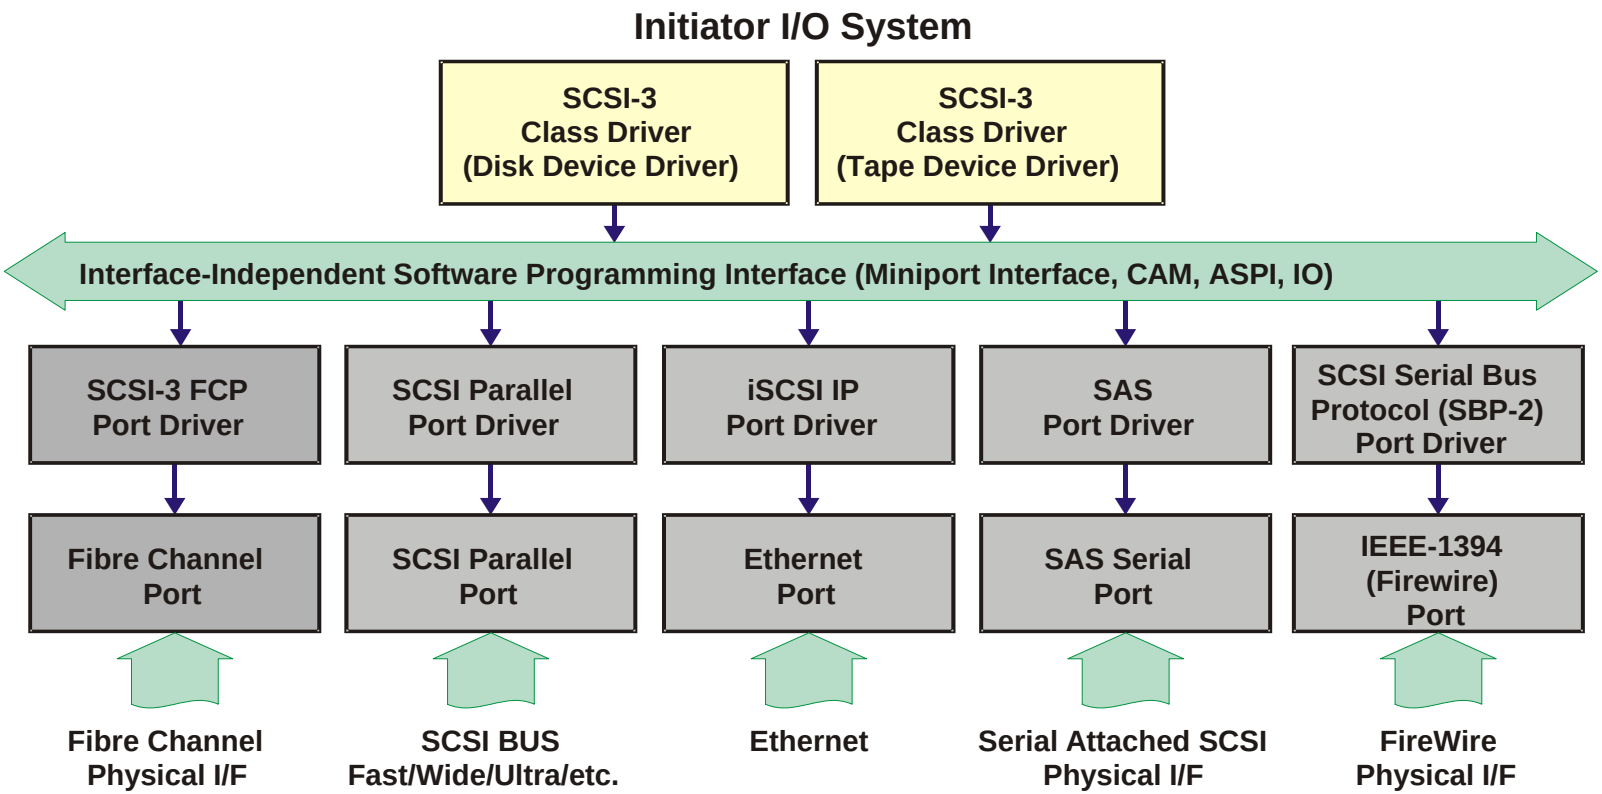

Protocol Service Model

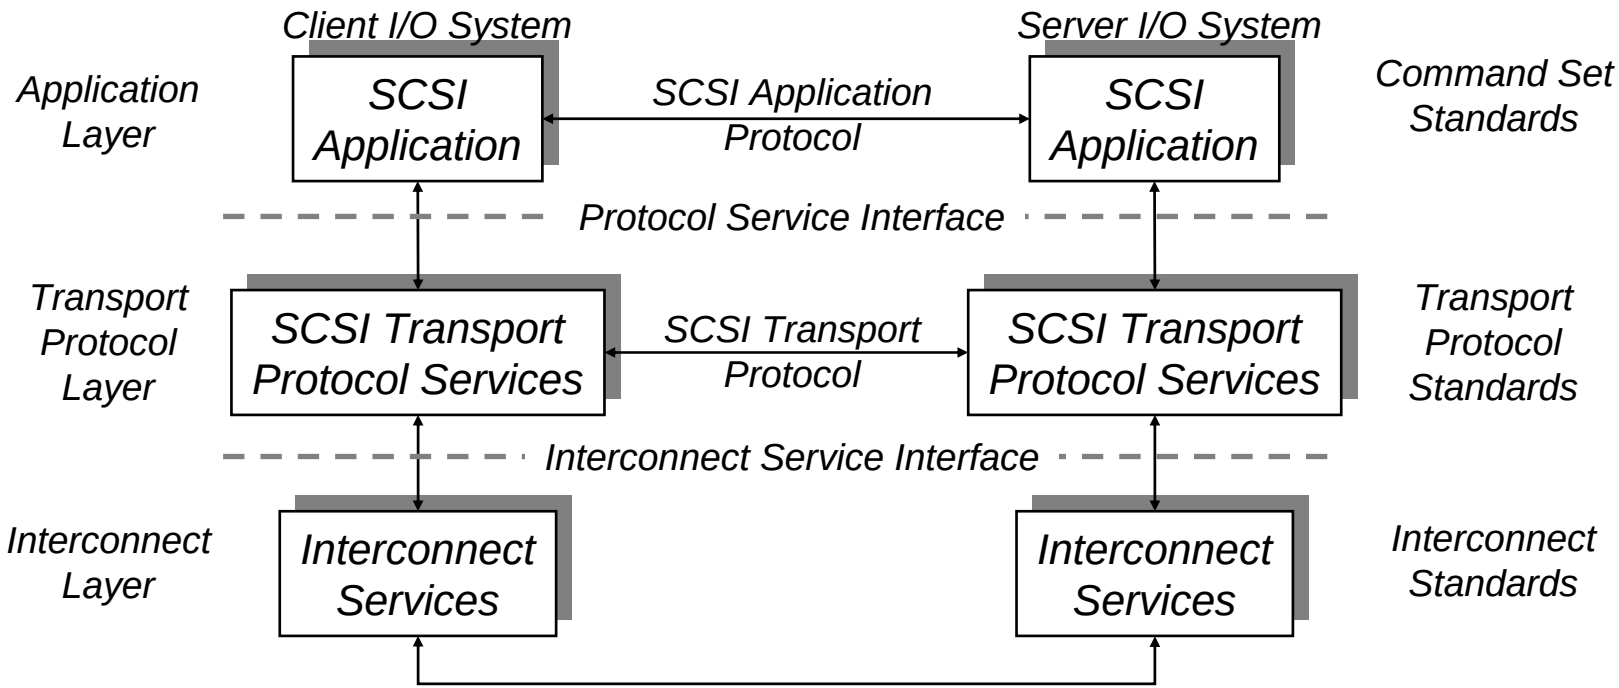

- Application layer: Clients and servers that originate and process SCSI I/O operations by means of a SCSI application protocol
- Transport protocol layer: Services and protocols through which clients and servers communicate
- Interconnect layer: Services, signaling mechanism and interconnect subsystem needed for the physical transfer of data from sender to receiver.

SCSI Command Sets

- The most appealing features behind the SCSI Architecture, i.e. the SCSI language
- Includes every type of storage device manufactured today; disk, tape, CD, storage enclosures, etc…
- Commands are manufacturer independent and are common for all types of devices
- SCSI device drivers and operating systems have utilized SCSI Command Sets since the late 80's
- Can be used in any storage architecture including Fibre Channel, iSCSI, InfiniBand, ATA/IDE, SATA (Serial ATA), SAS (Serial Attached SCSI), 1394 (Firewire)

SCSI Architectural Model

SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

Characteristics of the SCSI-3 Architecture

The SCSI Distributed Model

- SCSI is a client-server protocol.
- The *client* is called the ***Initiator*** (usually the OS I/O subsystem) and issues requests to the server.
- The *server* is called the ***Target*** (usually the SCSI controller that is part of a storage device) and receives, executes and returns Initiator requests and their associated responses.

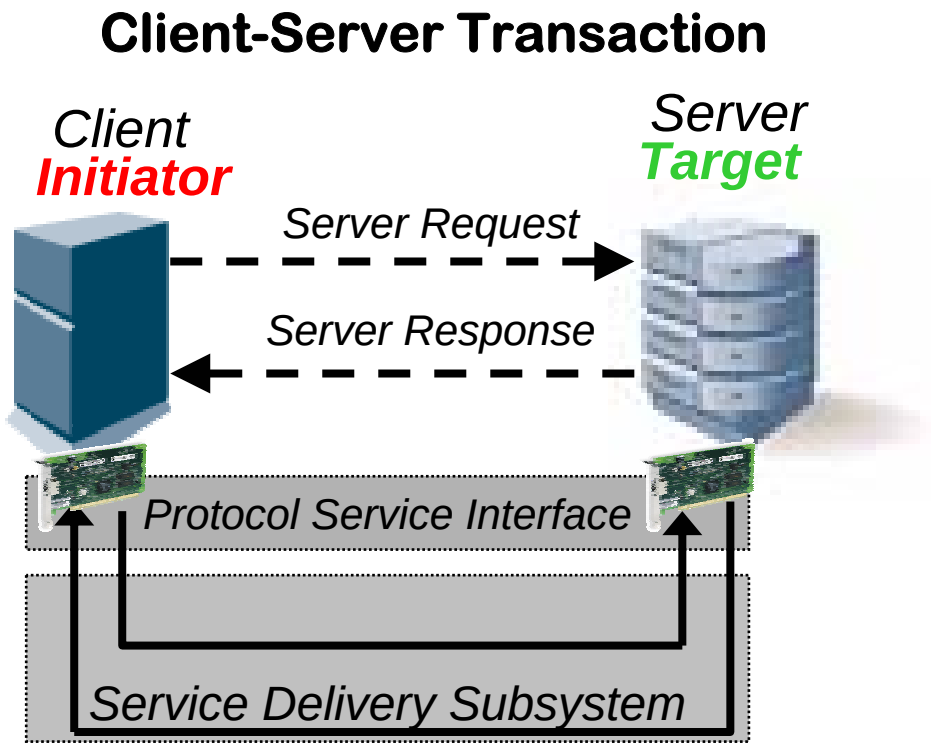

Client-Server Model

- A single Initiator can have multiple Application Clients.
- Targets have ONE Task Manager and one or more Logical Units (LU), which are numbered (LUN). The Task Manager:
  - **controls the sequencing of one or more tasks within a LU**
  - **carries out the task management functions**
  - **has the authority to modify Service Requests that have already been received by the target**
- The Device Server processes operations and directs them to a specific LUN.

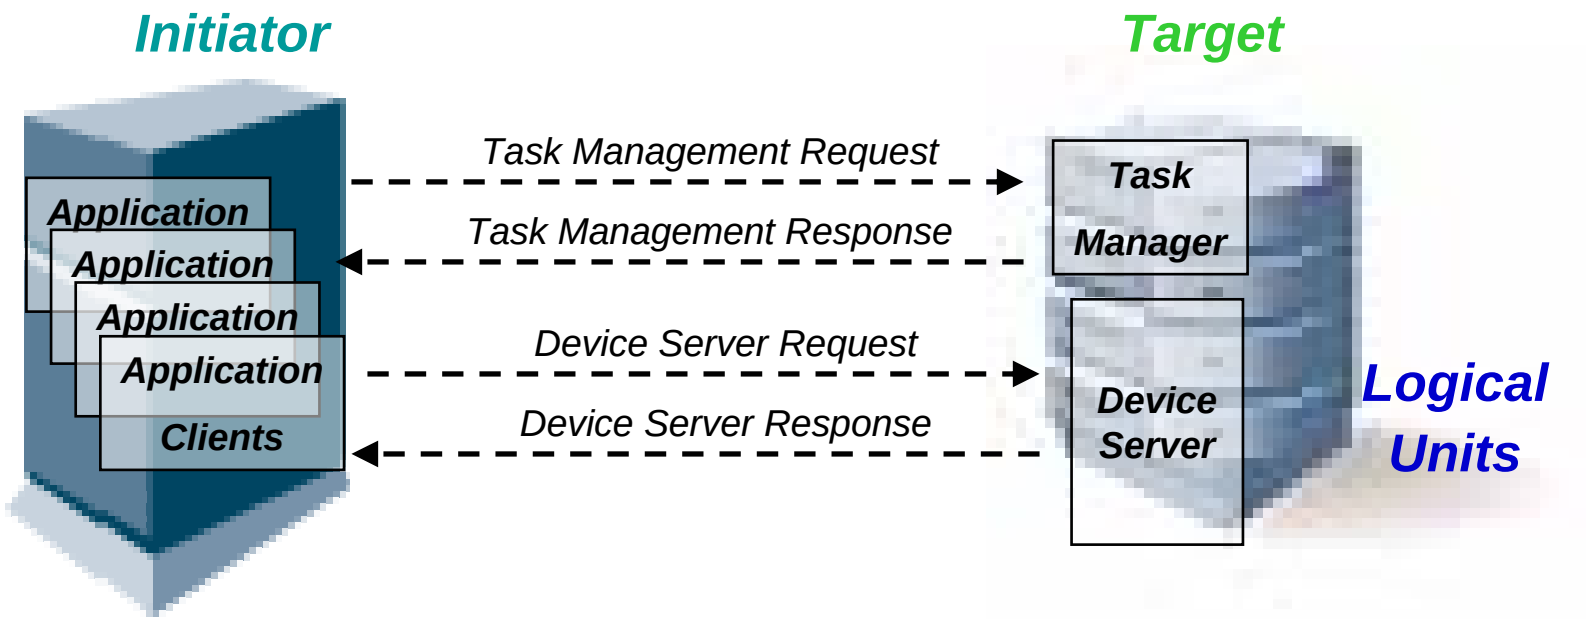

Re-cap SCSI Terminology

- SCSI is a *standard* that defines an interface between an ***Initiator*** (usually a computer) and a ***Target*** (usually a storage device such as a hard disk, tape backup, or storage array).
- ***Interface*** refers to connectors, cables, electrical signals, optical signals and the command protocol that allow initiators and targets to communicate.
- ***Logical Units*** are a subset of Target devices which can allow for scalability.
- The ***Queue*** (or Task Set) is used to hold pending Commands (Tasks) that the Target will execute

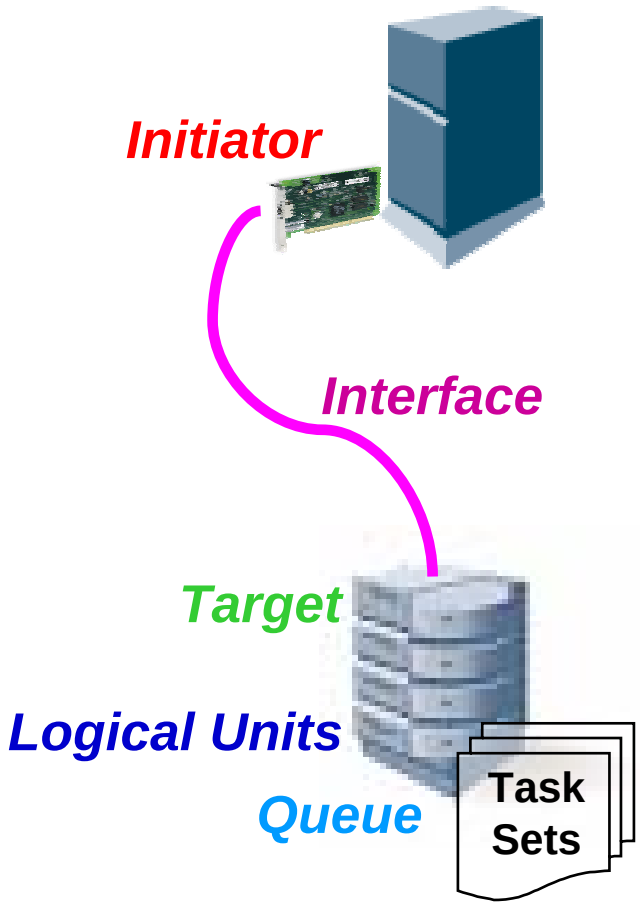

Task Attributes

- There are four types of Task attributes that can affect how the Target executes each task (Command)
  - **Simple**
    - **Target can execute in any order**
    - **Target will typically apply a performance algorithm to numerous simple tasks**
  - **Ordered**
    - **Target must execute all ordered tasks in the order they are received**
    - **Any task prior to ordered must be executed before ordered task**
  - **Head of queue**
    - **Informs Target to insert the task into the front of the queue**
  - **Auto Contingent Allegiance (ACA)**
    - **Used when the Target enters into an error condition for a command that has previously executed**

I/O Operations

SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

Characteristics of SCSI I/O Operations

I/O Operation Model

- There are two categories of Protocol Services:
  - **Execute Command and Confirmation Services**
  - **Data Transfer Services**
- There are three main phases of an I/O operation that includes a data transfer:
  1. **Command: Send required command and parameters via a Command Descriptor Block (CDB)**
  2. **Data: Transfer data in accordance with the command**
  3. **Status: Receive confirmation of command execution**

SCSI I/O Operations

- At a minimum, SCSI I/O Operations consist of:
  - **An Initiator issuing a SCSI Command**
  - **A Target returning completion Status**
  - **There is no "Data" transfer between Initiator and Target**
- Types of Commands that do not move Data
  - **Test Unit Ready**
  - **Start/Stop Unit**
  - **Rewind**

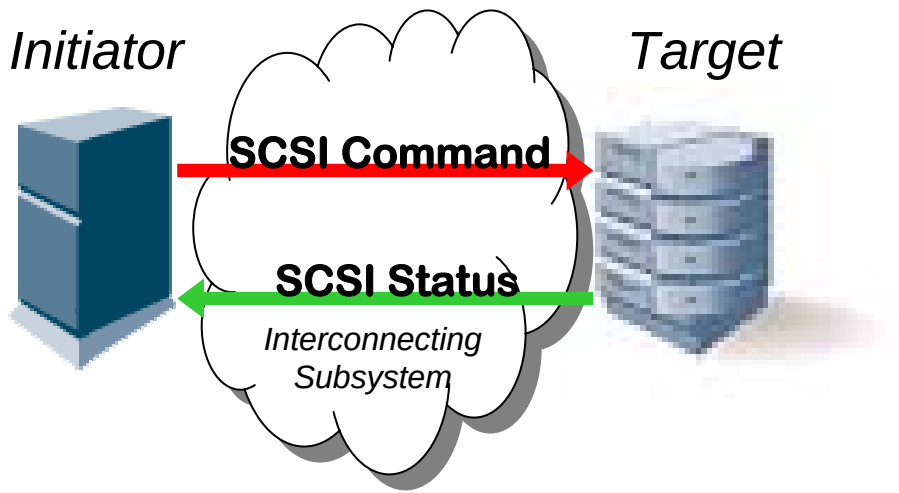

Operations that move Data

- When an Initiator and Target need to exchange information
  - **They utilize a Data phase**
  - **Data In transmits information from the Target to the Initiator**
  - **Data Out transmits information from the Initiator to the Target**
- Data can be transmitted all at once or take numerous Data phases to complete information transfers
- Types of Commands
  - **Read or Write**
  - **Inquiry**

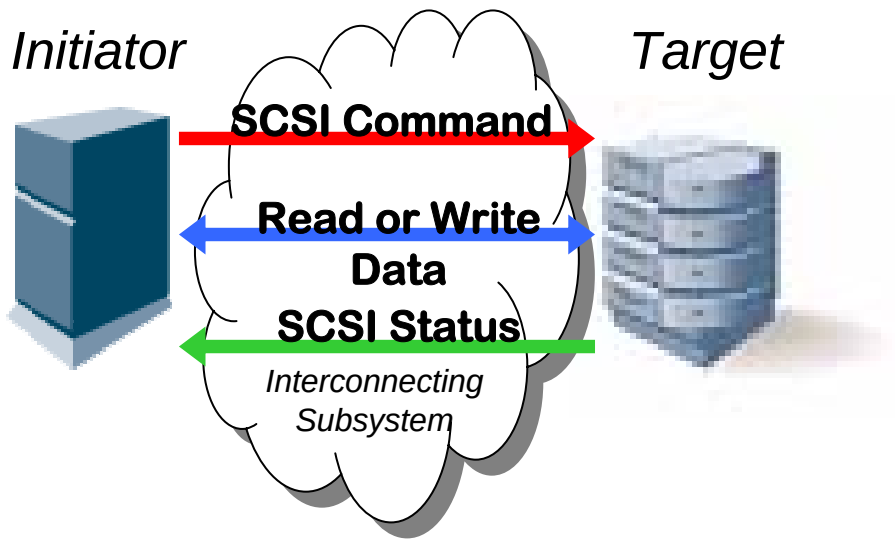

Command Descriptor Block

- A Command is executed by sending a Command Descriptor Block (CDB) to the Target
- For each CDB
  - **The first byte of the CDB is the Operation Code**
  - **The last byte of the CDB is the Control Byte**
  - **The format of the Operation Code and Control Byte is identical for every SCSI Command in the SCSI Universe**

Example Six Byte CDB

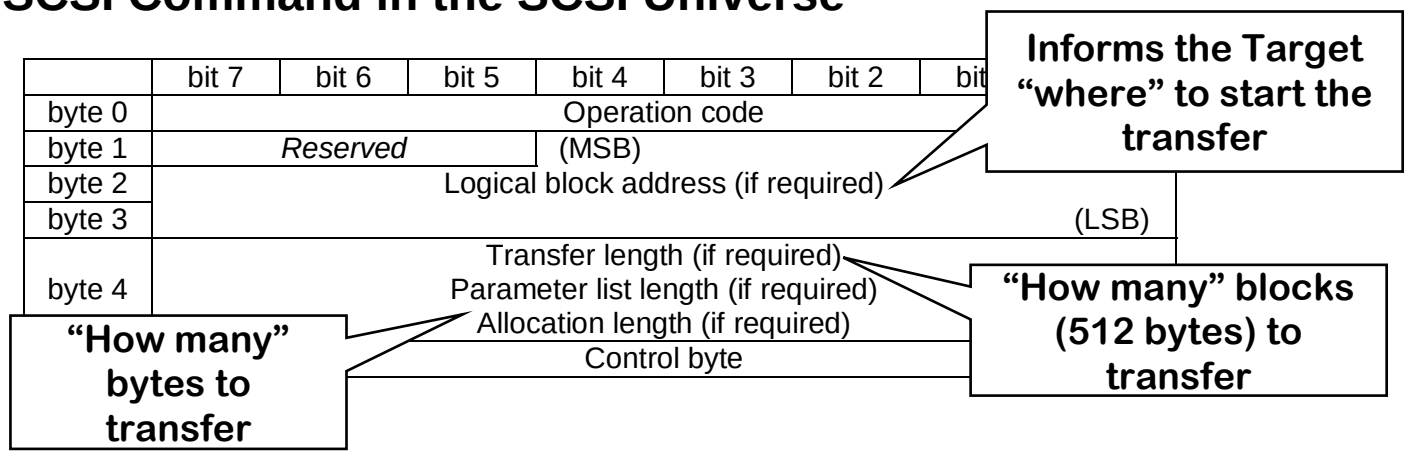

Other CDB Formats

CBDs can be:

- **10 bytes**
- **12 bytes**
- **16 bytes**
- **or even variable bytes in length**

Ten Byte CDB

	bit 7	bit 6	bit 5	bit 4	bit 3	bit 2	bit 1	bit 0
byte 0	Operation code							
byte 1	*Reserved*			*Service Action (if required)*				
byte 2	(MSB)							
byte 3	Logical block address (if required)							
byte 4								
byte 5								(LSB)
byte 6	Reserved							
byte 7	(MSB) Transfer length (if required) Parameter List Length (if required)							
byte 8	Allocation length (if required)							(LSB)
byte 9	Control byte							

Twelve Byte CDB

	bit 7	bit 6	bit 5	bit 4	bit 3	bit 2	bit 1	bit 0
byte 0	Operation code							
byte 1	*Reserved*			*Service Action (if required)*				
byte 2	(MSB)							
byte 3	Logical block address (if required)							
byte 4								
byte 5								(LSB)
byte 6	(MSB)							
byte 7	Transfer length (if required)							
byte 8	Parameter list length (if required)							
byte 9	Allocation length (if required)							(LSB)
byte 10	Reserved							
byte 11	Control byte							

SCSI Commands all Devices

Standard SCSI Commands:

- **Disk**
- **Tape**
- **Storage Enclosures**
- **Disk Arrays**
- **CD**
- **WORM**
- **Media Changers**
- **Every device imaginable or not yet created**

Op. Code	Command Name	Type
00h	TEST UNIT READY	M
03h	REQUEST SENSE	Z
12h	INQUIRY	M
15h	MODE SELECT(6)	Z
18h	COPY	O
1Ah	MODE SENSE(6)	Z
1Ch	RECEIVE DIAGNOSTIC RESULTS	O
1Dh	SEND DIAGNOSTIC	O
39h	COMPARE	O
3Ah	COPY AND VERIFY	O
3Bh	WRITE BUFFER	Z
3Ch	READ BUFFER	O
4Ch	LOG SELECT	O
4Dh	LOG SENSE	O
55h	MODE SELECT(10)	Z
5Ah	MODE SENSE(10)	Z

SCSI Status

- Any time a SCSI Command is sent to a Target
  - **The Initiator expects a completion Status**
  - **This status can reflect successful or unsuccessful completion of the command**
- The Status may indicate
  - **Busy or Not Ready**
  - **Error condition exists for another command**
  - **Targets task set is full**

Status	hex
GOOD	00
CHECK CONDITION	02
CONDITION MET	04
BUSY	08
INTERMEDIATE	10
INTERMEDIATE-CONDITION MET	14
RESERVATION CONFLICT	18
Obsolete ** *was* COMMAND TERMINATED	22
TASK SET FULL * was QUEUE FULL	28
ACA ACTIVE **	*30*
all other codes	rsvd

SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

Parallel SCSI

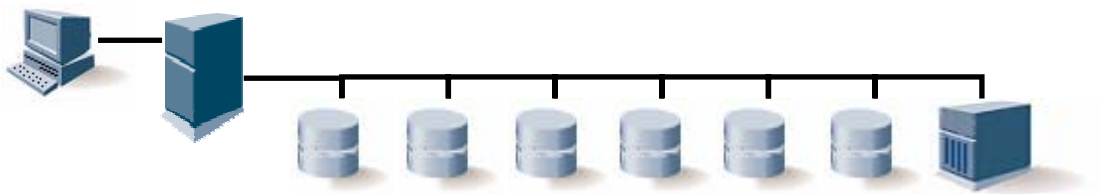

Parallel SCSI Characteristics

- Information can only go in one direction at a time
  - **Each piece (byte) of information is acknowledged**
  - **Information transfers are interlocked**
- A connection protocol is used before any information transfers
  - **Arbitration, Selection, and Message phases**
  - **Uses SCSI ID's to identify devices**
- Utilizes protocol "bus" phases to accomplish information transfers
  - **Command Phase to deliver the CDB**
  - **Data Phase to deliver customer data**
  - **Status Phase to deliver completion status**
- Can multi-task by disconnecting

Read Command Example

- Initiator connects to Target
  - **Once Initiator connects to Target, the Target is in control of the I/O Process**
- Initiator sends CDB information via Command Phases
- Target returns requested information via Data Phase
- Target returns completion information via Status Phase
- Target Disconnects via Message Phase

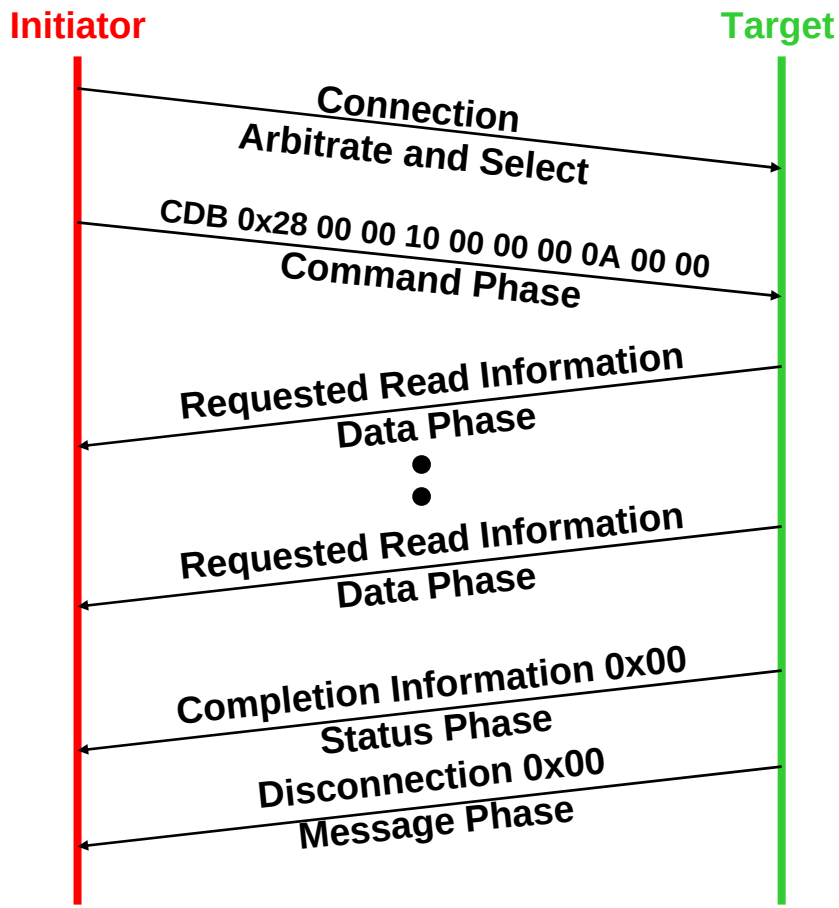

SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

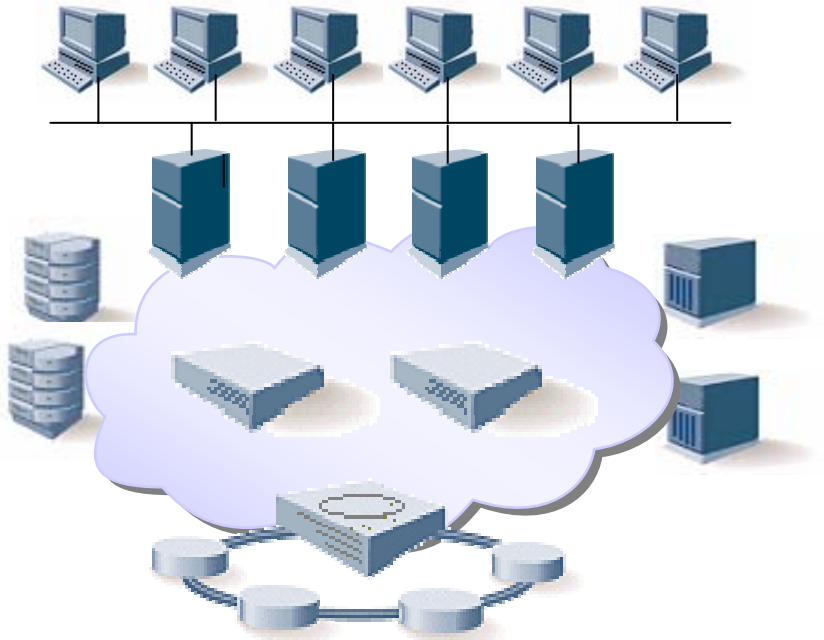

Fibre Channel

Fibre Channel Protocol Characteristics

- Fibre Channel creates frames to deliver SCSI Commands, Data, and Status information units (IU). FCP Frames include:
  - FCP_CMD frame
    - **Equivalent to parallel SCSI Command Phase**
  - FCP_XFER_RDY frame
    - **No actual parallel SCSI Equivalent**
  - FCP_DATA frame
    - **Equivalent to parallel SCSI Data Phase**
  - FCP_RSP frame
    - **Equivalent to parallel SCSI Status Phase**

Fibre Channel Frame

FCP_CMD IU
FCP_Data IU
FCP_XFER_RDY IU
FCP_RESP IU

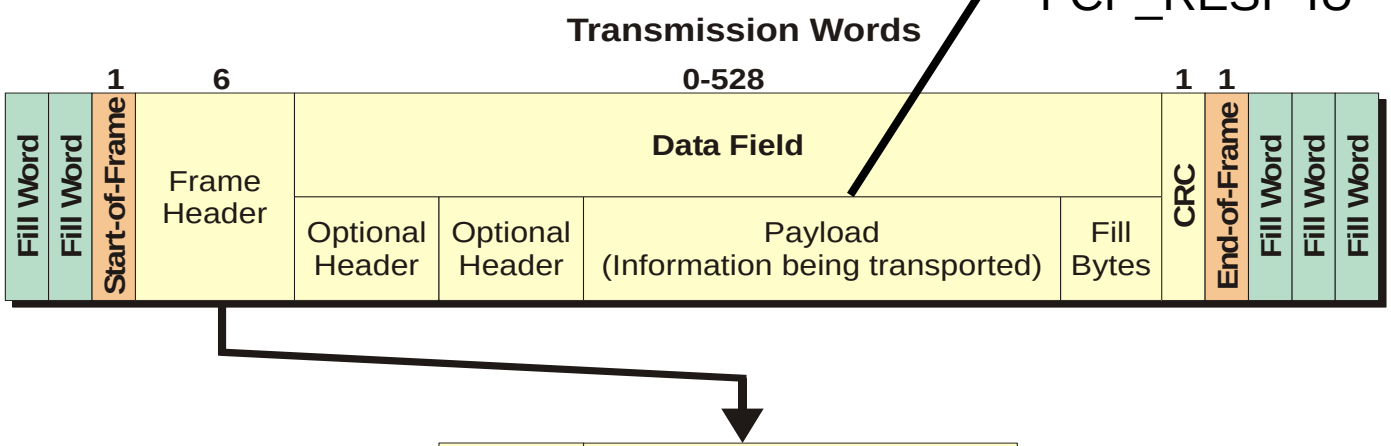

Frame type and content/function	R_CTL	Destination Address (D_ID)	*Where frame is being sent to*
Class-specific control information	CS_CTL	Source Address (S_ID)	*Where the frame came from*
Protocol Type in this frame	TYPE	Frame Control (F_CTL)	*Frame Control field*
Sequence this frame belongs to	SEQ_ID	DF_CTL / SEQ_CNT	*Sequential count of frames*
Originator Exchange ID	OX_ID	RX_ID	*Responder Exchange ID*
Multi-purpose parameter field	Parameter Field (PARM)		

Write Command Example

- Before devices can communicate in FC they must establish a session via login protocol
- Originator sends Write command to responder
- Responder replies with a ready to receive response
- Originator sends Data frames
- When all Data is transferred, Responder sends Status to Originator
- Session could end at this point but typically doesn't

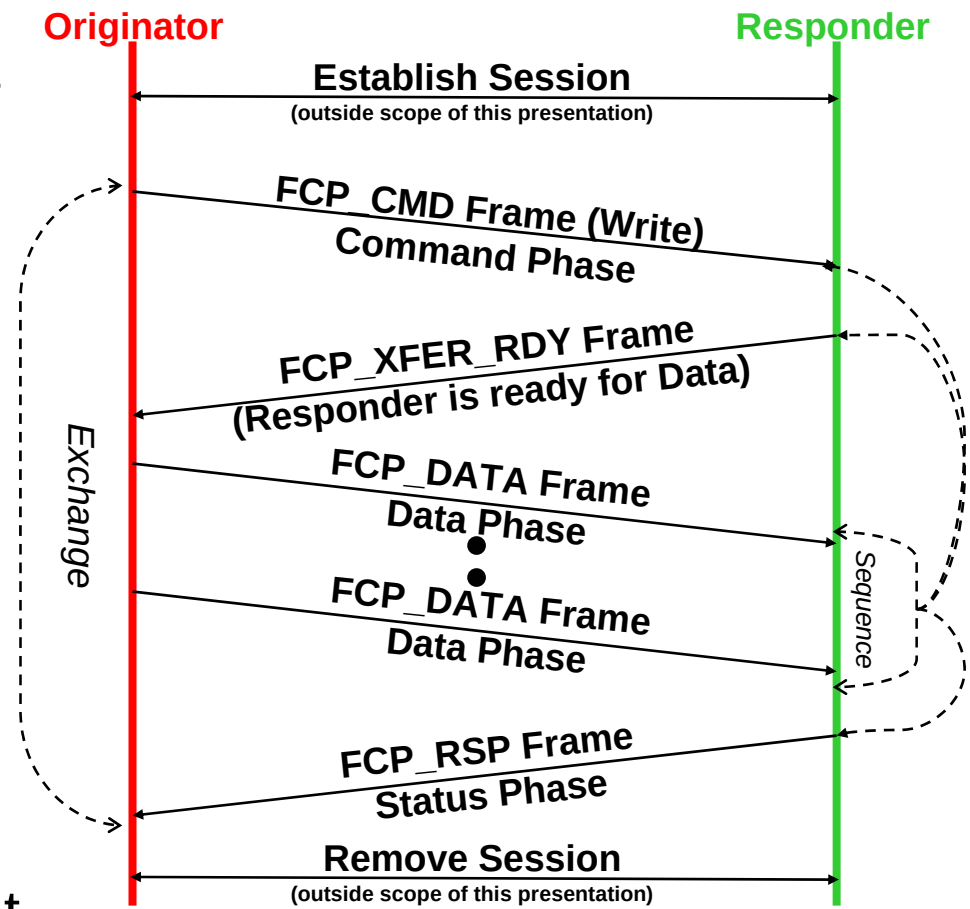

SNIA

STORAGE NETWORKING INDUSTRY ASSOCIATION

EDUCATION

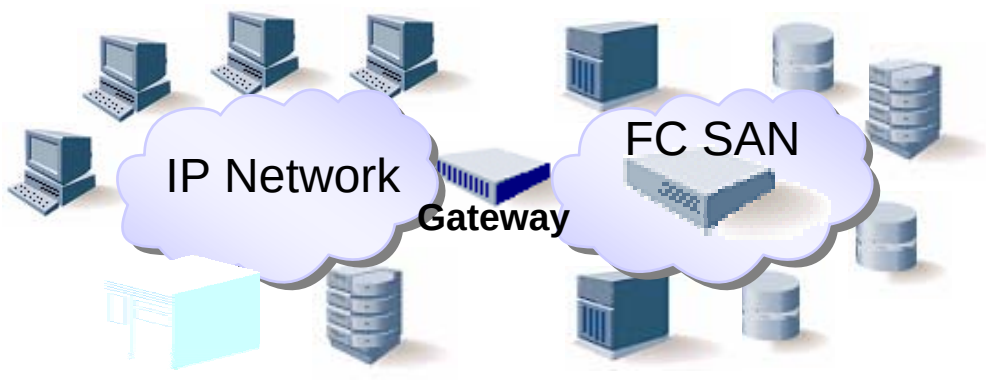

iSCSI

iSCSI Characteristics

- iSCSI is a SCSI transport protocol for mapping of block oriented storage data over TCP/IP networks.
- Storage generally identified for this application include disk, tape arrays and tape libraries.
- IP networks most applicable for this purpose are Gigabit Ethernet and in the future 10 Gigabit Ethernet, however for low performance applications 10/100bT will work.
- The iSCSI layer encapsulates the SCSI CDB into a iSCSI Protocol Data Unit (PDU) and forwards it to the Transmission Control Protocol (TCP) layer.
- The communications between the Initiator and Target will occur over one or more TCP connections.
- The TCP connections form a session and will carry the iSCSI PDU's. The sessions are given an ID called a Connection ID (CID). There are two parts of the ID, Initiator Session ID (ISID) and Target ID (TSID) and together make up an "I_T nexus".

iSCSI Encapsulation

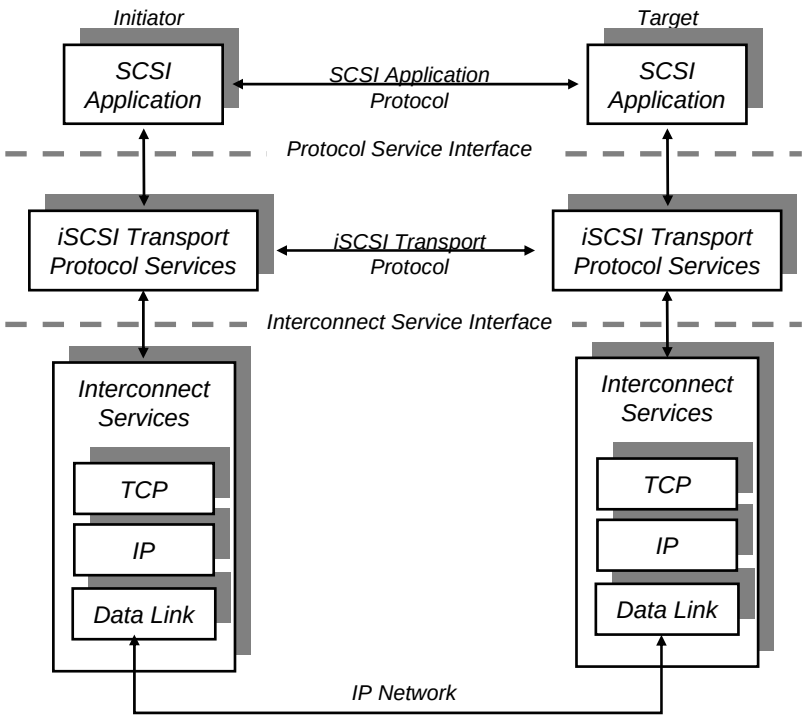

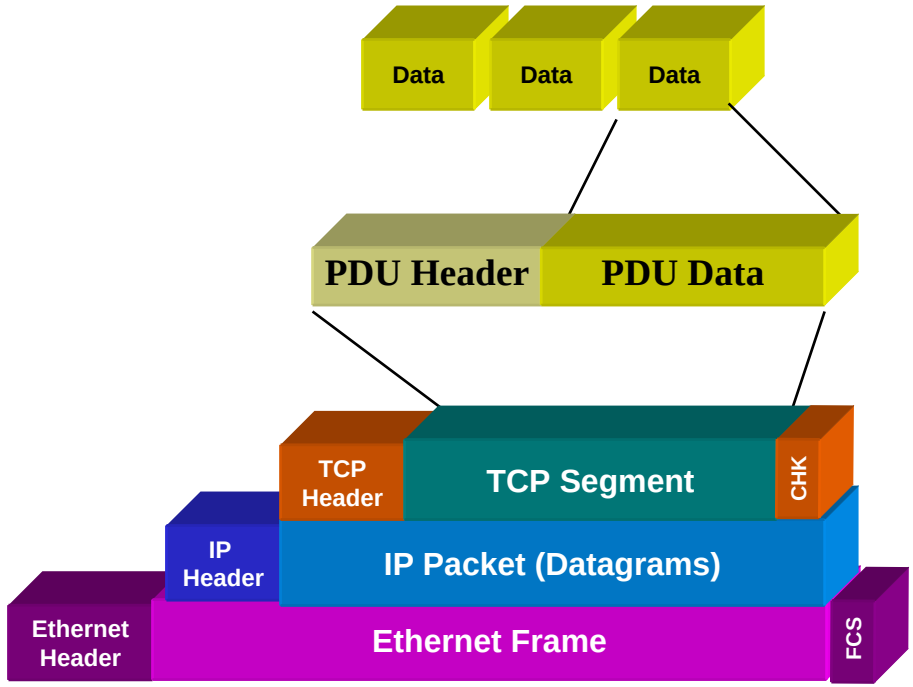

Write Command Example

- Before devices can communicate in iSCSI they must establish a session via login protocol
- Initiator sends Write command to Target
- Target replies with a ready to transmit R2T response
- Initiator sends Data frames
- When all Data is transferred, Target sends Status to Initiator
- Session could end at this point but typically doesn't

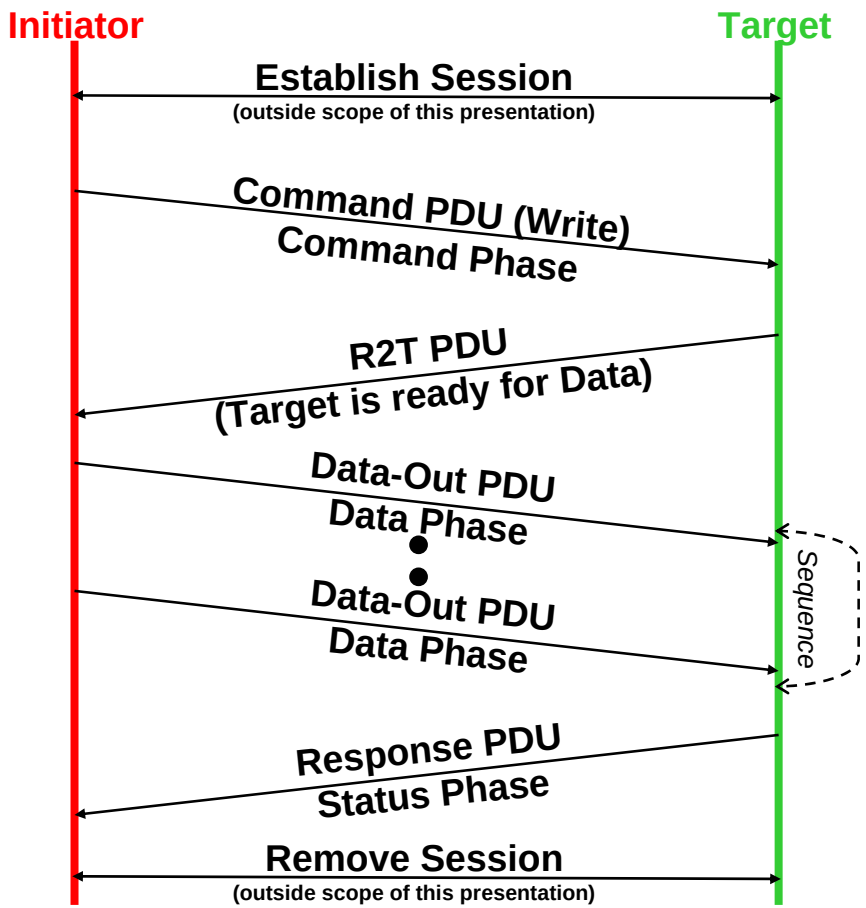

Read Command Example

- Before devices can communicate in iSCSI they must establish a session via login protocol
- Initiator sends Read command to Target
- Target sends Data frames
- When all Data is transferred, Target sends Status to Initiator
- Session could end at this point but typically doesn't

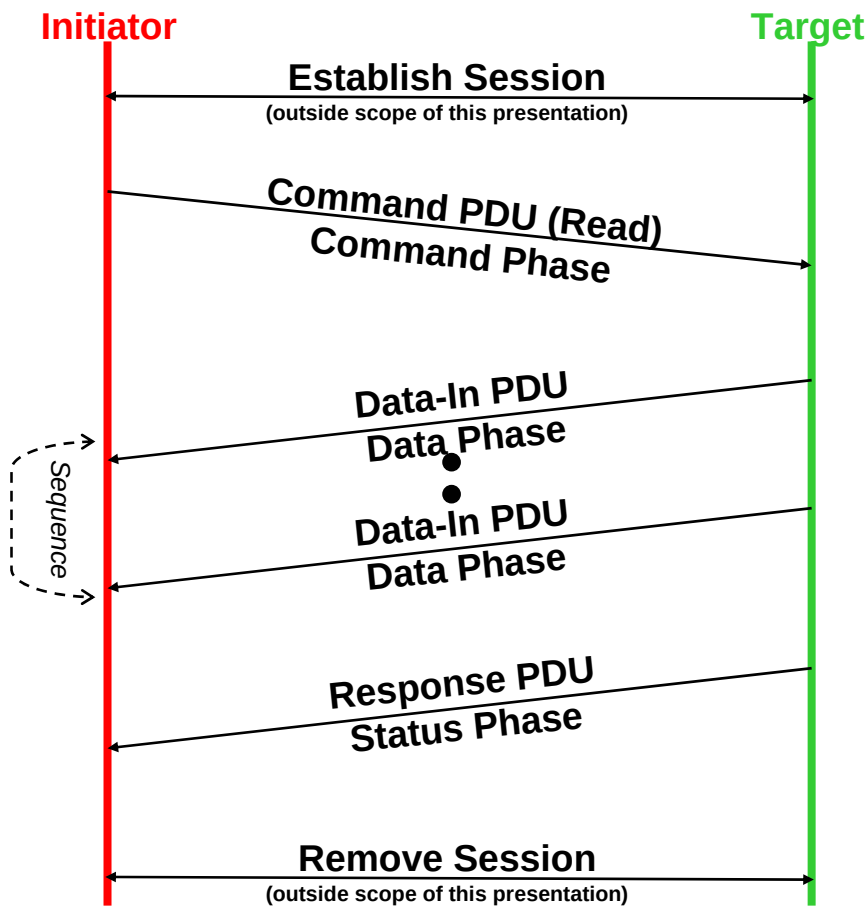

Summary

- SCSI is the language of LOVE
- Every storage architecture utilizes the SCSI model
- SCSI Commands sets are available for every type of storage device even including host-to-host
- An I/O Process is interface independent and is made up of:
  - Command Phase
  - Optional Data Phase
  - and a Status Phase
- SCSI language is here to stay and will be used in every storage architecture including:
  - SATA
  - SAS
- For more information on standards see: incits.org, t10.org, t11.org, t13.org, ietf.org, scsita.org, fibrechannel.org

Q&A / Feedback

- Please send any questions or comments on this presentation to SNIA: track-storage@snia.org

Many thanks to the following individuals for their contributions to this tutorial.

SNIA Education Committee

David Deming
Elaine Silber